

Introduction to Machine Learning

資工系 許涵義 109550200

Final Project – Report

1. Environment Details

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import csv
import math

from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.impute import SimpleImputer, KNNImputer
from sklearn.pipeline import make_pipeline
from sklearn.metrics import roc_auc_score, accuracy_score
from sklearn.model_selection import cross_validate, StratifiedKFold

from feature_engine.encoding import WOEEncoder

from colorama import Fore, Back, Style

from joblib import dump
from joblib import load

from google.colab import drive
drive.mount('/content/drive')
```

```
[ ] # Requirements #
# !pip install -r requirements.txt
# or
!pip install numpy pandas matplotlib scikit-learn joblib feature_engine colorama
```

Below are all the libraries needed to run the code, you can install the libraries by pip install libraries or simply install the requirements.txt.

2. Implementation Details

a. Train.ipynb

```
[ ] def impute_null_values(df, features, pcodes):
    # Impute null values in a dataframe using a KNNImputer model #

    # Create KNNImputer model with 15 neighbors
    model = KNNImputer(n_neighbors=15)
    # model = KNNImputer(n_neighbors=3)
    # model = KNNImputer(n_neighbors=5)

    for pcode in pcodes:
        # Get rows of data where product_code is equal to the current value of 'pcode'
        mask = df['product_code'] == pcode
        feature_subset = df.loc[mask, features]

        # Calculate the number of null values in the feature subset before imputation
        null_before = feature_subset.isnull().sum().sum()

        # Use the KNN model to impute null values in the feature subset
        imputed = model.fit_transform(feature_subset)

        # Calculate the number of null values in the imputed feature subset
        null_after = pd.isnull(imputed).sum().sum()

        # Calculate the number of null values that were imputed
        null_imputed = null_before - null_after

        # Print a message indicating how many null values were imputed for the current product code
        print(f"Imputing Product Code {pcode}: {null_imputed} null values imputed")

        # Replace the feature subset in the original data with the imputed version
        df.loc[mask, features] = imputed

    return df

[ ] train = impute_null_values(train, features, train['product_code'].unique())
# train.query("m3_missing != 0 and m5_missing != 0").head()
```

```
1 # Initialize the train Path #
# myPath = "tabular-playground-series-mq-2022"
myPath = "/content/drive/myDrive/ML/Final-Project/tabular-playground-series-mq-2022"
train = pd.read_csv(f'{myPath}/train.csv')

# For Feature Selection #
train2 = pd.DataFrame.copy(train)
if 'failure' in train2:
    target = train2.pop('failure')
print(f'train {train.shape}')
train.head()
```

id	product_code	loading	attribute_0	attribute_1	attribute_2	attribute_3	measure
0	0	A	80.10	material_7	material_8	9	5
1	1	A	84.89	material_7	material_8	9	5
2	2	A	82.43	material_7	material_8	9	5
3	3	A	101.07	material_7	material_8	9	5
4	4	A	108.06	material_7	material_8	9	5

5 rows x 26 columns

```
2 def scale_function(X, cols):
    scaler = StandardScaler()
    # scaler = MinMaxScaler()

    X_scaled = scaler.fit_transform(X[cols])
    X_scaled = pd.DataFrame(X_scaled, columns=cols, index=X.index)
    X_scaled = pd.concat([X.drop(cols, axis=1), X_scaled], axis=1)

    assert len(X) == len(X_scaled)

    return X_scaled
```

As there are null values present in both the training and test datasets, so we need to perform imputation (replacing null values) on the data. One method is to use *KNNImputer* from *sklearn.impute* module. By using the *impute_null_values* function we can impute all the null values and assign it by *KNNImputer*. After testing with different imputer module and the parameter of nearest neighbor, *KNNImputer* with *n_neighbor* = 15 has the highest accuracy for the model.

As we need to perform feature selection later, scaling the data is an important preprocessing step. Scaling the data ensures that all features are on a similar scale, thus making it easier to compare the importance of different feature. The `scale_function` will be used when training the model.

i. Training without Feature Selection

```
def train_model(select_feature, X, y):
    # Initialize lists to store feature importances
    importance_list = []

    # Initialize variables to store average AUC and accuracy scores
    avg_auc = 0
    avg_acc = 0

    # Initialize variables to store OOF AUC and accuracy scores
    oof_auc = 0
    oof_acc = 0

    # Initialize empty arrays to store OOF predictions
    oof_preds_proba = np.zeros(len(X))
    oof_preds = np.zeros(len(X))

    # Define stratified k-fold cross-validation object with 5 splits
    kf = StratifiedKFold(n_splits=5, shuffle=True, random_state=0)

    # Loop through the k-fold splits
    for fold_idx, (train_idx, val_idx) in enumerate(kf.split(X, y)):

        # Split data into training and validation sets
        X_train, X_val = X.iloc[train_idx], X.iloc[val_idx]
        y_train, y_val = y.iloc[train_idx], y.iloc[val_idx]

        # Scale data using the scale_data() function
        X_train_scaled = scale_function(X_train, select_feature)
        X_val_scaled = scale_function(X_val, select_feature)

        # Testing HyperParameters
        # Fit logistic regression model on training data
        # model = LogisticRegression(max_iter=1000, C=0.0001, penalty='l2', solver='newton-cg')
        # model = LogisticRegression(C=0.1, penalty='l2', solver='newton-cg', random_state=0)
        model = LogisticRegression(penalty='l1', C=0.01, solver='liblinear', random_state=1)
        model.fit(X_train_scaled[select_feature], y_train)

        # Make predictions on validation data
        val_preds_proba = model.predict_proba(X_val_scaled[select_feature])[0, 1]
        val_preds = model.predict(X_val_scaled[select_feature])

        # Compute and accumulate AUC and accuracy scores for validation set
        score = roc_auc_score(y_val, val_preds_proba)
        avg_auc += roc_auc_score(y_val, val_preds_proba) / 5
        avg_acc += accuracy_score(y_val, val_preds) / 5

        # Store OOF predictions
        oof_preds_proba[val_idx] = val_preds_proba
        oof_preds[val_idx] = val_preds
```

```
# Append feature importances to lists
importance_list.append(model.coef_.ravel())

# Print current fold number
print(f"Fold {fold_idx}: ROC-AUC = {score:.5f}")

# Print average AUC and accuracy scores
print(f"{Fore.GREEN}{Style.BRIGHT}Average auc = {round(avg_auc, 5)}")

# Compute and print OOF AUC and accuracy scores
oof_auc = roc_auc_score(y, oof_preds_proba)
oof_acc = accuracy_score(y, oof_preds)
print(f"{Fore.BLUE}{Style.BRIGHT}OOF auc = {round(oof_auc, 5)}")

# Create dataframe of feature importances
importance_df = pd.DataFrame(np.array(importance_list).T, index=X[select_feature].columns)
importance_df['mean'] = importance_df.mean(axis=1).abs()
importance_df['feature'] = X[select_feature].columns
importance_df = importance_df.sort_values('mean', ascending=False).reset_index().head(10)

# Plot top 20 features as horizontal bar chart
plt.figure(figsize=(14,4))
plt.barh(importance_df.index, importance_df['mean'], color='lightgreen')
plt.gca().invert_yaxis()
plt.xticks(ticks=importance_df.index, labels=importance_df['feature'])
plt.title('Logistic Regression Feature Importances')
plt.show()

# Return model
return model
```

2.1 Training without Feature Selection

```
# Initial Training without any Feature Selection
X = train.drop(['failure'], axis=1)
y = train['failure'].astype(int)

print(features)
model1 = train_model(features, X, y)

['loading', 'measurement_0', 'measurement_1', 'measurement_2', 'measurement_3', 'measurement_4',
Fold 0: ROC-AUC = 0.66180
Fold 1: ROC-AUC = 0.59172
Fold 2: ROC-AUC = 0.58052
Fold 3: ROC-AUC = 0.59276
Fold 4: ROC-AUC = 0.58383
Average auc = 0.58997
OOF auc = 0.58989]
```


This `train_model` function is designed to train a logistic regression model on a provided dataset using k-fold cross-validation. The script first splits the data into training and validation sets using a stratified k-fold cross-validation object with 5 splits. This ensures that each fold contains roughly the same proportions of the different classes in the dataset.

It applies feature scaling to the data by using a `scale_function` which scales the data according to the feature selected. A logistic regression model is then fit to the training data with the selected features. After several trials and error with the hyperparameter of the Logistic Regression model, using the `'liblinear'` solver and L1 regularization with a *regularization strength* of 0.01, `random_state=1` obtained the best result.

Predictions are then made on the validation data, and the area under the receiver operating characteristic (ROC) curve (AUC) and accuracy scores are calculated and accumulated over the 5 folds. The average AUC and accuracy scores are then

printed at the end of the script. Additionally, it also stores out-of-fold predictions, which are used to compute out-of-fold AUC and accuracy scores.

Furthermore, the *train_model* function also appends feature importances to lists and creates a dataframe of feature importances sorted by the mean importance of the feature, it prints the top 10 feature importances.



ii. Training with Feature Selection

Now, to improve the model accuracy, we implement Feature Engineering.

```

2.2 Training with Feature Selection

[19] train2['m3_missing'] = train2['measurement_3'].isnull().astype(np.ints)
train2['m5_missing'] = train2['measurement_5'].isnull().astype(np.ints)

X_train2 = impute_null_values(train2, features, train2['product_code'].unique())

Imputing Product Code A:....
-> 3849 null values imputed
Imputing Product Code B:....
-> 3975 null values imputed
Imputing Product Code C:....
-> 4344 null values imputed
Imputing Product Code D:....
-> 3973 null values imputed
Imputing Product Code E:....
-> 4132 null values imputed

[21] # Encoding for attribute_0 using WoEEncoder #
woe_encoder = WoEEncoder(variables=['attribute_0'])
woe_encoder.fit(X_train2, target)

X_train2 = woe_encoder.transform(X_train2)
# X_test2 = woe_encoder.transform(X_test2)

print(X_train2.columns)
model2 = train_model(optimized_features, X_train2, target)

Index(['id', 'product_code', 'loading', 'attribute_0', 'attribute_1',
       'attribute_2', 'attribute_3', 'measurement_0', 'measurement_1',
       'measurement_2', 'measurement_3', 'measurement_4', 'measurement_5',
       'measurement_6', 'measurement_7', 'measurement_8', 'measurement_9',
       'measurement_10', 'measurement_11', 'measurement_12', 'measurement_13',
       'measurement_14', 'measurement_15', 'measurement_16', 'measurement_17',
       'm3_missing', 'm5_missing', 'measurement(3*5)', 'missing(3*5)', 'area'],
      dtype='object')
Fold 0: ROC-AUC = 0.60033
Fold 1: ROC-AUC = 0.59179
Fold 2: ROC-AUC = 0.58138
Fold 3: ROC-AUC = 0.59186
Fold 4: ROC-AUC = 0.58482
Average auc = 0.59003
OOF auc = 0.58998

[22] # Important Feature Engineering #
X_train2['measurement(3*5)'] = X_train2['measurement_3'] * X_train2['measurement_5']
X_train2['missing(3*5)'] = X_train2['m5_missing'] * X_train2['m3_missing']
X_train2['area'] = X_train2['attribute_2'] * X_train2['attribute_3']


X_train2[['area', 'm3_missing', 'm5_missing', 'missing(3*5)', 'measurement_3', 'measurement_5', 'measurement(3*5)']].query('m3_missing != 0 and m5_missing != 0').head()

[23] optimized_features = [
    'loading',
    'attribute_0',
    'area',
    'measurement_17',
    'measurement_0',
    'measurement_1',
    'measurement_2',
    'm3_missing',
    'm5_missing',
    'measurement(3*5)',
    'missing(3*5)'
]

```

The above code utilizes feature engineering to select a subset of the available features for training the logistic regression model. The selected features are listed in the *optimized_features* list and include various measurements, attributes, and calculated values. The process of selecting these features likely involved analyzing the feature importance of the model trained with all available features, as well as researching and implementing additional features based on analysis of feature correlations and discussion on Kaggle. Breakdown of the features:

1. Analyzing the feature importance from training the model with all features, we pick the top features loading, measurement 0 - 2 (integer distribution), and measurement 17(float distribution) among the other features.
2. Researching a lot of kaggle discussions about feature correlations and EDA analysis, we create the features:
 - *attribute_0*
 - *area*
 - *m3_missing*
 - *m5_missing*
 - *measurement(3*5)*
 - *missing(3*5)*
 - a) *attribute_0* contains string value that represent categorical features that should be one-hot-encoded
 - b) *area* is obtained by multiplying attribute_2 and attribute_3 with the assumption that these data might represent the dimensions (width * length)
 - c) *m3_missing* and *m5_missing* is based on this discussion <https://www.kaggle.com/competitions/tabular-playground-series-aug-2022/discussion/342319>
 - d) *measurement(3*5)* and *missing(3*5)* is based on this discussion <https://www.kaggle.com/competitions/tabular-playground-series-aug-2022/discussion/343368>



b. Inference.ipynb

```
1. Data Preprocessing for Test Dataset
We also need to implement impute_null_values and scale_function to test dataset

def impute_null_values(df, features, pcodes): ...

def scale_function(X, cols): ...

test2['m0_missing'] = test2['measurement_3'].isnull().astype(np.inta)
test2['m5_missing'] = test2['measurement_5'].isnull().astype(np.inta)
X_test2 = impute_null_values(test2, features, test2['product_code'].unique())

Inputting Product Code B:....
-> 4873 null values inputed
Inputting Product Code G:....
-> 3924 null values inputed
Inputting Product Code H:....
-> 3883 null values inputed
Inputting Product Code I:....
-> 3829 null values inputed

[130] # Need X_train2 for fitting #
train2['m0_missing'] = train2['measurement_3'].isnull().astype(np.inta)
train2['m5_missing'] = train2['measurement_5'].isnull().astype(np.inta)
X_train2 = impute_null_values(train2, features, train2['product_code'].unique())

Inputting Product Code A:....
-> 3449 null values inputed
Inputting Product Code B:....
-> 3975 null values inputed
Inputting Product Code C:....
-> 4344 null values inputed
Inputting Product Code D:....
-> 3973 null values inputed
Inputting Product Code E:....
-> 4132 null values inputed

[131] # encoding for attribute_0 using hotencoder #
woe_encoder = woeEncoder(variables=['attribute_0'])
woe_encoder.fit(X_train2, target)

# X_train2 = woe_encoder.transform(X_train2)
X_test2 = woe_encoder.transform(X_test2)

[133] # Important Feature Engineering #
X_test2['measurement(3%)'] = X_test2['measurement_3'] * X_test2['measurement_5']
X_test2['missing(3%)'] = X_test2['m0_missing'] * X_test2['m5_missing']
X_test2['area'] = X_test2['attribute_2'] * X_test2['attribute_3']

X_test2[['area', 'm0_missing', 'm5_missing', 'missing(3%)', 'measurement_3', 'measurement_5',
        'area_m0_missing_m5_missing_missing(3%) measurement_3 measurement_5 measurement(3%)']]

area m0_missing m5_missing missing(3%) measurement_3 measurement_5 measurement(3%)
29 24 1 1 1 17.684200 17.152000 301.947236
1931 24 1 1 1 17.936333 17.179333 308.134249
2656 24 1 1 1 18.138400 17.224267 312.420638
3036 24 1 1 1 17.800867 17.098133 304.361592
5917 63 1 1 1 17.651933 17.196733 303.555580

[134] optimized_features = [
    'loading',
    'attribute_0',
    'area',
    'measurement_1%',
    'measurement_0',
    'measurement_1',
    'measurement_2',
    'm0_missing',
    'm5_missing',
    'measurement(3%)',
    'missing(3%)'
]

[135] # scaling #
X_test2_scaled = scale_function(X_test2, optimized_features)
```

```
2. Load the pretrained Model

[136] predmodel2 = LogisticRegression(penalty='l1', C=0.01, solver='liblinear', random_state=1)
# predmodel2 = LogisticRegression(max_iter=200, C=0.0001, penalty='l2', solver='newton-cg')

modelPATH = '/content/drive/MyDrive/ML-Final-Project'
predmodel2 = load(f'{modelPATH}/my_best_model3.joblib')
y_pred2 = predmodel2.predict_proba(X_test2_scaled[optimized_features])[0,1]

3. Writing to Submission.csv

[137] X_test2 = X_test2.reset_index().set_index('id').drop('index', axis=1)

submission = pd.DataFrame({'id': X_test2.index, 'failure': y_pred2})
submission.to_csv('109550200_v2.csv', index=False)
submission.head()
```

	id	failure
0	26570	0.198423
1	26571	0.180607
2	26572	0.190013
3	26573	0.191483
4	26574	0.332543

In order to use the trained model to make predictions, the test dataset must first be preprocessed to ensure that it is in the same format as the training data. This includes imputing any null values, selecting the same features as those used in the training data, and scaling the data using the selected features.

Then, load the pre-trained model from a file, which has been trained and saved from the previous process, in this case *my_best_model3.joblib*. Use the loaded model to predict the test data, it will output the predicted classes or probabilities of the test dataset. Finally writing the prediction to the submission.csv

3. Final Result

Github Link : <https://github.com/NicoA07/MachineLearning-FinalProject>

Model Weight Link: https://github.com/NicoA07/MachineLearning-FinalProject/blob/main/my_best_model3.joblib

Playground Prediction Competition

Tabular Playground Series - Aug 2022

Practice your ML skills on this approachable dataset!

Kaggle - 1,888 teams - 4 months ago

Overview Data Code Discussion Leaderboard Rules Team Submissions **Late Submission**

Submissions

You selected 0 of 2 submissions to be evaluated for your final leaderboard score. Since you selected less than 2 submission, Kaggle auto-selected up to 2 submissions from among your public best-scoring unselected submissions for evaluation. The evaluated submission with the best Private Score is used for your final score.

Submissions evaluated for final score

All Successful Selected Errors

Recent

Submission and Description	Private Score	Public Score	Selected
109550200_v2.csv Complete (after deadline) · now	0.59108	0.58524	<input type="checkbox"/>
109550200_v1.csv Complete (after deadline) · 1s ago	0.58989	0.58682	<input type="checkbox"/>

Reference:

<https://www.kaggle.com/competitions/tabular-playground-series-aug-2022/discussion/349299>
<https://www.kaggle.com/competitions/tabular-playground-series-aug-2022/discussion/343960>
<https://www.kaggle.com/competitions/tabular-playground-series-aug-2022/discussion/349541>
<https://www.kaggle.com/competitions/tabular-playground-series-aug-2022/discussion/342126>
<https://www.kaggle.com/competitions/tabular-playground-series-aug-2022/discussion/342319>
<https://www.kaggle.com/competitions/tabular-playground-series-aug-2022/discussion/343368>